

Lecture 08: Sep 17, 2018

Linear Regression

- *SLR*
- *MLR*
- *Factors and Design Matrices*

James Balamuta
STAT 385 @ UIUC

Announcements

- **hw03** is due on **Friday, Sep 21st, 2018 @ 6:00 PM**
- **Quiz 04** covers Week 3 contents @ [CBTF](#).
 - Window: Sep 18th - 20th
 - Sign up: <https://cbtf.engr.illinois.edu/sched>
- **hw01** grade reports released on GitHub.
 - Post on forum detailing how to interpret the [grade reports](#).
 - Got caught using GitHub's web interface? Let's chat.

Last Time

- **Data Structures**

- $1D$, $2D$, and n Dimensions
- Homogenous (Same) vs. Heterogenous (Different/Mixed)

- **Coercion**

- Changing data from one form to the another either implicitly (R) or explicitly (You).

- **Missingness and NULL**

- The lack of recorded data vs. the lack of an object being created.

Lecture Objectives

- **Fitting** a linear regression model with data
- Constructing **Design Matrices** associated with MLR
- Benefits associated with **using factors**.

SLR

Previously


Understand the Algorithm

... weeks of programming saved hours of planning ...

- *What* logic is being used?
- *How* does the logic apply in a procedural form?
- *Why* is this logic present?

SLR in Practice

... finding the line of best fit between two variables ...

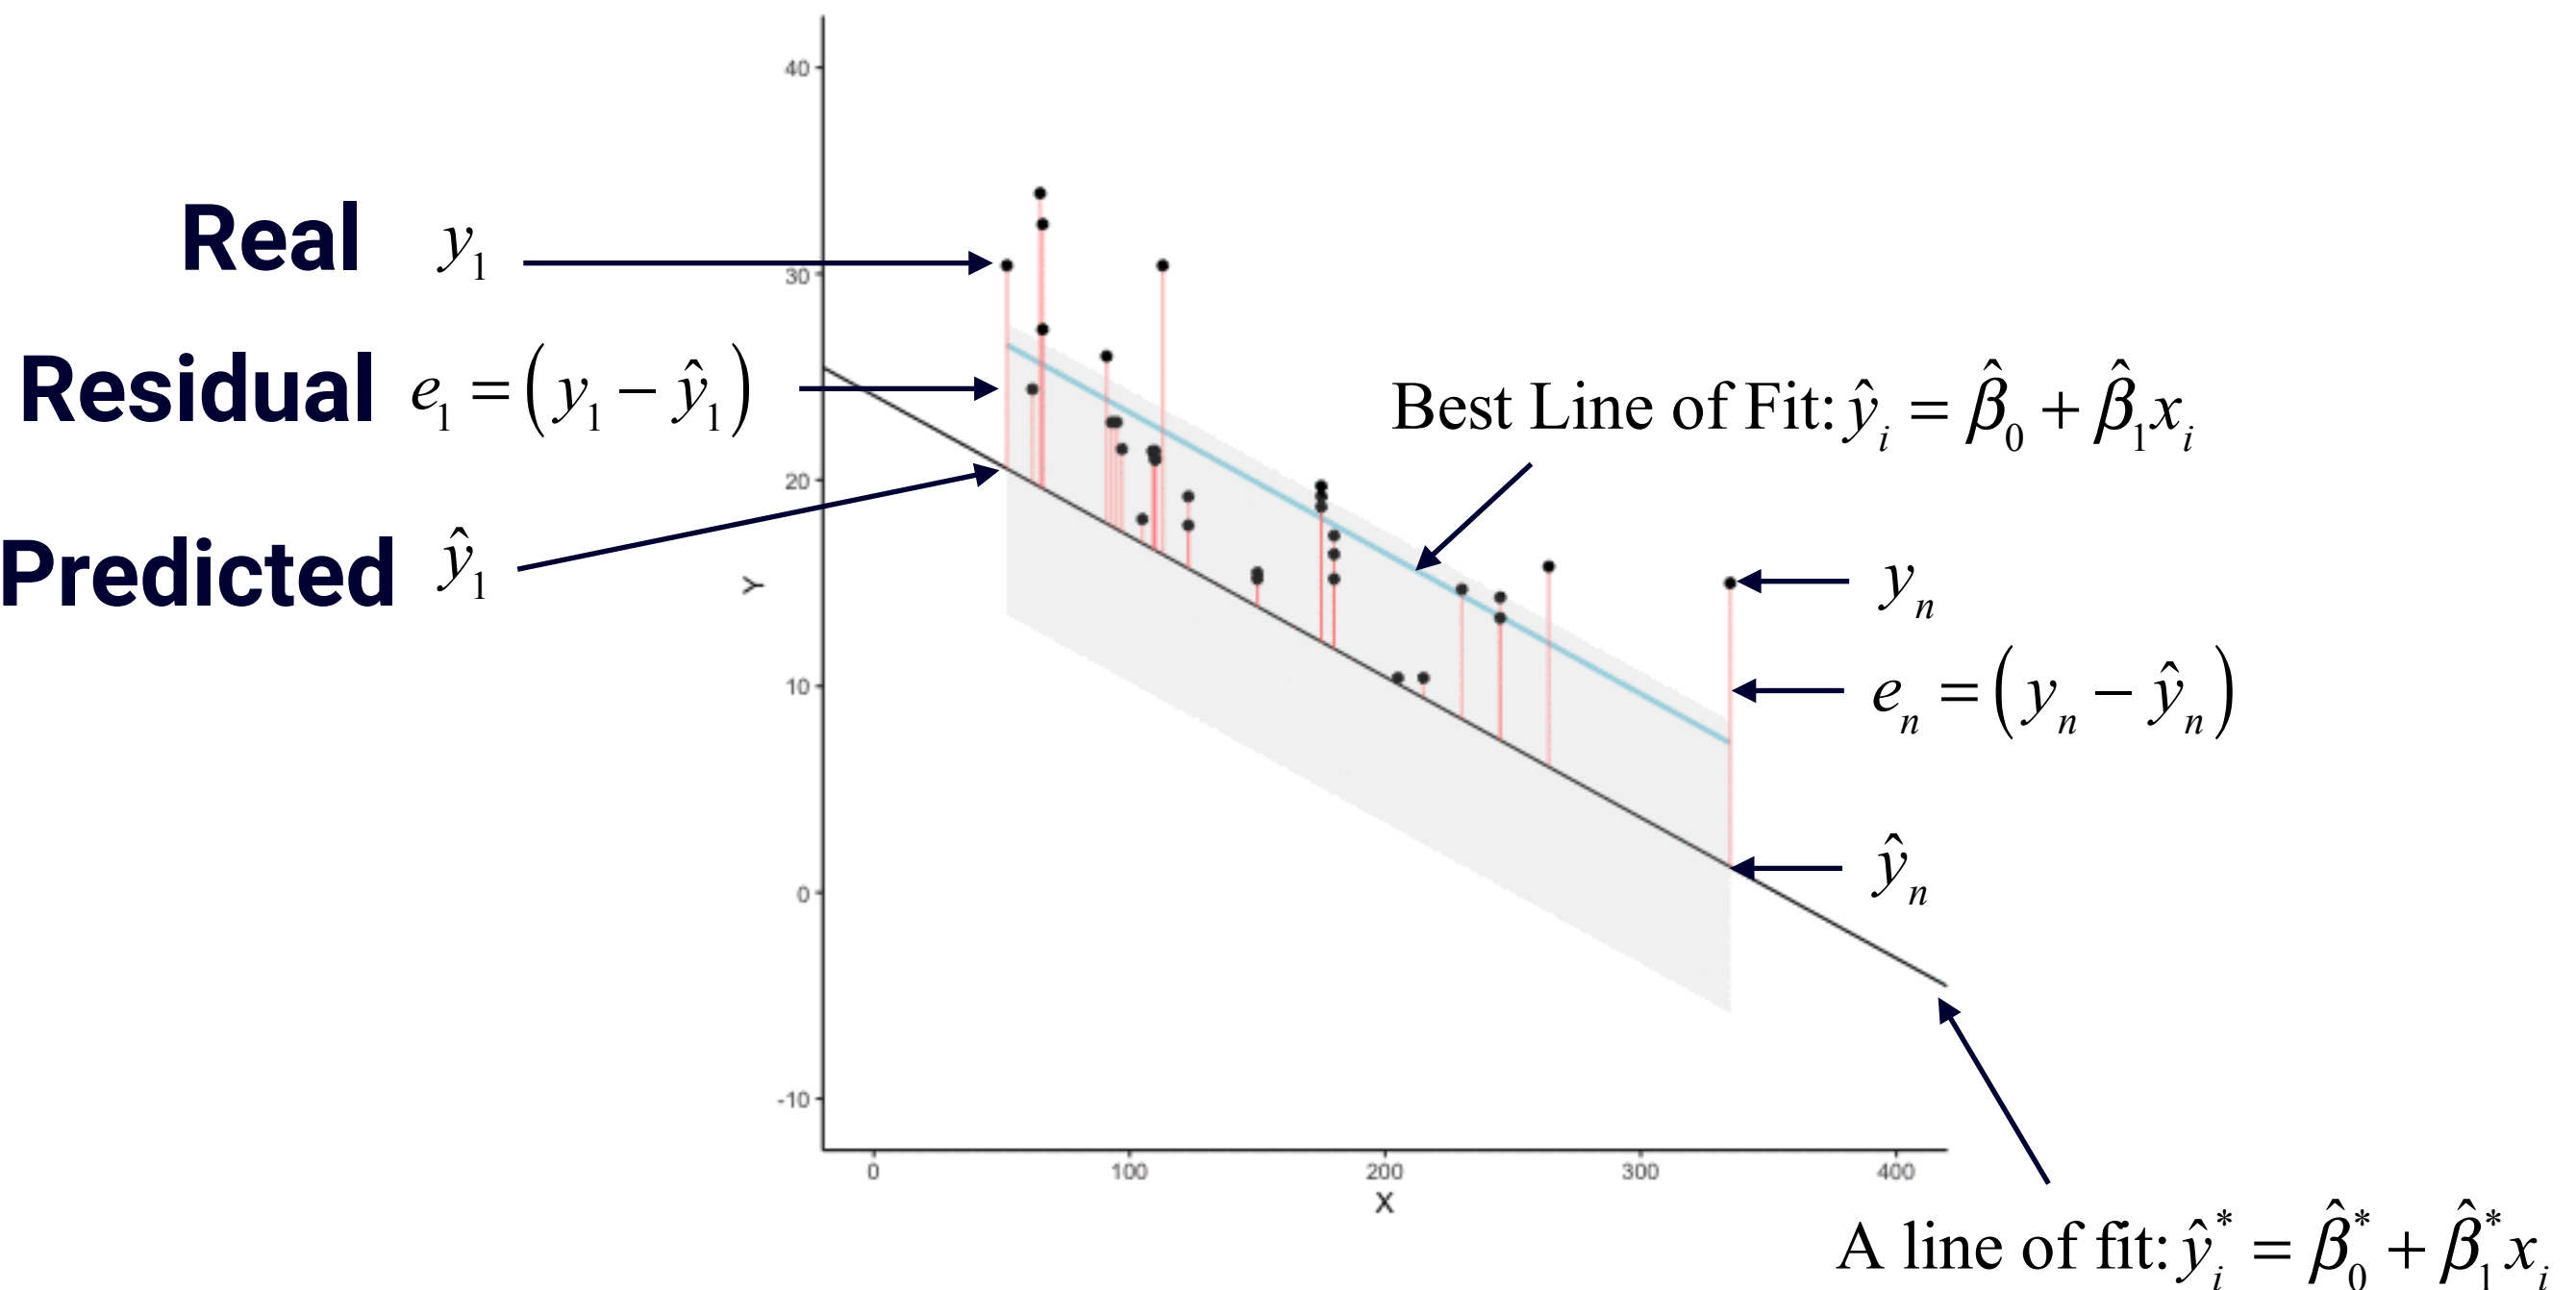


[Source](#)

-
- * The **blue** line represents the optimal line of best fit.
 - ** The **black** line represents the current line of fit.
 - *** The **red** lines represent distance from points. The goal is to *minimize* these values.

SLR Labeled

... graph components ...



Simple Linear Regression

... handling two parameters ...

Scalar-form

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$$

Expand Matrix

$$\underbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}_{\text{Responses}} = \underbrace{\begin{pmatrix} 1 & x_{1,1} \\ \vdots & \vdots \\ 1 & x_{n,1} \end{pmatrix}}_{\text{Design Matrix}} \underbrace{\begin{pmatrix} \beta_0 \\ \beta_1 \end{pmatrix}}_{\text{Parameters}} + \underbrace{\begin{pmatrix} \varepsilon_1 \\ \vdots \\ \varepsilon_n \end{pmatrix}}_{\text{Error}}$$

Matrix-form

$$Y_{n \times 1} = X_{n \times 2} \beta_{2 \times 1} + \varepsilon_{n \times 1}$$

n is the number of observations, there are 2 variables, \mathbf{X} provides the design matrix for the variables, \mathbf{y} is response vector, β is the parameter or coefficient vector and ε is the random error vector

Estimating Parameters

... goal is to **minimize the Residual Sum Squared** (*RSS/red lines*) ...


$$\hat{\beta} = \arg \min_{\beta_0, \beta_1 \in \mathbb{R}} \sum_{i=1}^n \left(y_i - \underbrace{(\beta_0 + \beta_1 x_i)}_{\hat{y}_i} \right)^2$$

**Analytical
Solutions**

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = \frac{S_{xy}}{S_{xx}}$$


$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \quad \bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$



Simulating Data for SLR

... generating a test data set ...

Generated Data
with 500 observations



```
# Set seed for reproducibility  
set.seed(9123)
```

```
# Number of observations  
n = 500
```

```
# Generate a single predictor  
x_i = rnorm(n, mean = 2)
```

```
# Design matrix w/ intercept:  $n \times 2$   
X = cbind(1, x_i)
```


```
# True beta values:  $2 \times 1$   
beta = c(2.5, 4)
```

```
# Response variable with error:  $n \times 1$   
y = X[, 1] * beta[1] +  
    X[, 2] * beta[2] +  
    rnorm(n, sd = 2) # error term
```

Analytical SLR Solution

... solving using equations ...

Estimated Coefficients on Generated Data
with 500 observations



```
# Mean of Response (y)
y_mu = mean(y)
```

```
# Mean of Predictor (x)
x = X[, 2]
x_mu = mean(x)
```

```
# Estimate Slope
beta1_hat =
  sum((x - x_mu) * (y - y_mu)) /
  sum((x - x_mu) ^ 2)
```

```
# Estimate Intercept
beta0_hat =
  y_mu - beta1_hat * x_mu
```

```
# Check Parameter Difference
cbind(c(beta0_hat, beta1_hat), beta)
#      beta_hat  beta
# [1,] 2.511657  2.5
# [2,] 4.027985  4.0
```

Optimizing RSS

... minimizing point distance ...


```
# Write the cost function to minimize
min_rss_slr = function(par, X, y) {
  rss = sum((y - (X[,1]*par[1] + X[, 2] * par[2]))^2)
  return(rss)
}
```

```
# Initial beta parameters values
beta_init = c(0, 0)
```

```
# Perform the minimization
model_opt = optim(par = beta_init,
  fn = min_rss_slr, method = "BFGS",
  control = list(trace = TRUE),
  X = X, y = y)
```

```
# initial value 4545.000693
# final value 4542.177432
# converged
```

```
# Check parameter difference
cbind(model_opt$par, beta)
#      beta_hat beta
# [1,] 2.511657 2.5
# [2,] 4.027985 4.0
```



MLR

Multiple Linear Regression

... scaling SLR to multiple predictors ...

Scalar-form

$$y_i = \beta_0 + \beta_1 x_{i,1} + \beta_2 x_{i,2} + \cdots + \beta_{p-1} x_{i,p-1} + \varepsilon_i$$

Expand Matrix

$$\underbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}_{n \times 1 \text{ Responses}} = \underbrace{\begin{pmatrix} 1 & x_{1,1} & \cdots & x_{1,p-1} \\ \vdots & \vdots & & \vdots \\ 1 & x_{n,1} & \cdots & x_{n,p-1} \end{pmatrix}}_{n \times p \text{ Design Matrix}} \underbrace{\begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_{p-1} \end{pmatrix}}_{p \times 1 \text{ Parameters}} + \underbrace{\begin{pmatrix} \varepsilon_1 \\ \vdots \\ \varepsilon_n \end{pmatrix}}_{n \times 1 \text{ Error}}$$

Matrix-form

$$Y_{n \times 1} = X_{n \times p} \beta_{p \times 1} + \varepsilon_{n \times 1}$$

n is the number of observations, **p is the number of variables**, X is where predictors are stored, \mathbf{y} is response vector, β is the parameter or coefficient vector and ε is the random error vector

SLR vs MLR

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

$$Y_{n \times 1} = X_{n \times 2} \hat{\beta}_{2 \times 1}$$

$$\underbrace{\begin{pmatrix} \hat{y}_1 \\ \vdots \\ \hat{y}_n \end{pmatrix}}_{n \times 1 \text{ Responses}} = \underbrace{\begin{pmatrix} 1 & x_{1,1} \\ \vdots & \vdots \\ 1 & x_{n,1} \end{pmatrix}}_{n \times 2 \text{ Design Matrix}} \underbrace{\begin{pmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \end{pmatrix}}_{2 \times 1 \text{ Parameters}}$$

$$\hat{\beta} = \arg \min_{\beta_0, \beta_1 \in \mathbb{R}} \sum_{i=1}^n \left(y_i - (\beta_0 + \beta_1 x_i) \right)^2$$

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{i,1} + \hat{\beta}_2 x_{i,2} + \cdots + \hat{\beta}_{p-1} x_{i,p-1}$$

$$\hat{Y}_{n \times 1} = X_{n \times p} \hat{\beta}_{p \times 1}$$

$$\underbrace{\begin{pmatrix} \hat{y}_1 \\ \vdots \\ \hat{y}_n \end{pmatrix}}_{n \times 1 \text{ Responses}} = \underbrace{\begin{pmatrix} 1 & x_{1,1} & \cdots & x_{1,p-1} \\ \vdots & \vdots & & \vdots \\ 1 & x_{n,1} & \cdots & x_{n,p-1} \end{pmatrix}}_{n \times p \text{ Design Matrix}} \underbrace{\begin{pmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \\ \vdots \\ \hat{\beta}_{p-1} \end{pmatrix}}_{p \times 1 \text{ Parameters}}$$

$$\hat{\beta} = \arg \min_{\beta \in \mathbb{R}^p} \left\| \mathbf{y}_{n \times 1} - X_{n \times p} \beta_{p \times 1} \right\|^2$$

Derivations of MLR

<http://thecoatlessprofessor.com/statistics/multiple-linear-regression-proofs/>

STAT 420 and STAT 425 will focus more on the derivations.

Fitting MLR

... using linear regression ...

Formula

The underlying model for the data

Data

Data to regress over

```
model_fit = lm(y ~ x1 + x2, data = data_set)
```

Inside the *formula*, model terms can be combined using **+** to give:

$$y_i = \beta_0 + \beta_1 x_{i,1} + \beta_2 x_{i,2} + \epsilon_i$$

* β_0 represents the intercept term and is automatically included.

** The intercept can be removed by subtracting 1, e.g. $y \sim x - 1$

Example MLR Fit

... fitting data to mtcars ...

```
# Construct a multiple linear regression (3 variables: wt, qsec, intercept)
```

```
model_fit = lm(mpg ~ wt + qsec, data = mtcars)
```

```
# View the estimated parameters
```

```
model_fit
```

```
# Call:
```

```
# lm(formula = mpg ~ wt + qsec, data = mtcars)
```

```
#
```

```
# Coefficients:
```

```
# (Intercept)      wt      qsec
```

```
# 19.7462    -5.0480     0.9292
```

How formulas generate a design matrix

Constructing a design matrix **with** the intercept

$X = \text{model.matrix}(\text{mpg} \sim \text{wt} + \text{qsec}, \text{data} = \text{mtcars})$

Building a design matrix **without** the intercept

$X_{\text{noint}} = \text{model.matrix}(\text{mpg} \sim \text{wt} + \text{qsec} - \mathbf{1}, \text{data} = \text{mtcars})$

Intercept

(Intercept)	wt	qsec
1	2.62	16.46
1	2.875	17.02
1	2.32	18.61

X

No Intercept

wt	qsec
2.62	16.46
2.875	17.02
2.32	18.61

X_{noint}

Inference

... understanding the regression model ...

```
summary(lm(mpg ~ wt + qsec, data = mtcars))
```

$$\mathbf{e} = \mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}$$

$$e_i = y_i - \left(\sum_{j=0}^{p-1} x_{i,j} \hat{\beta}_j \right)$$

Call:

```
lm(formula = mpg ~ wt + qsec, data = mtcars)
```

Residuals:

```
summary(residuals(model))
```

Min	1Q	Median	3Q	Max
-4.3962	-2.1431	-0.2129	1.4915	5.7486

Coefficients:

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$s.e.(\hat{\beta}_j) = \sqrt{\text{cov}(\hat{\beta}_j)} = \sqrt{\sigma_{RSS}^2 (\mathbf{X}^T \mathbf{X})_{jj}^{-1}}$$

$$\hat{\sigma}_{RSS} = \sqrt{\frac{\mathbf{e}^T \mathbf{e}}{n-p}} = \sqrt{\frac{1}{n-p} \sum_{i=1}^n e_i^2}$$

$$R^2 = 1 - \frac{RSS}{TSS}$$

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	19.7462	5.2521	3.760	0.000765 ***
wt	-5.0480	0.4840	-10.430	2.52e-11 ***
qsec	0.9292	0.2650	3.506	0.001500 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.596 on 29 degrees of freedom

Multiple R-squared: 0.8264, Adjusted R-squared: 0.8144

F-statistic: 69.03 on 2 and 29 DF, p-value: 9.395e-12

$$t_{\hat{\beta}} = \frac{\hat{\beta} - \beta_0}{s.e.(\hat{\beta})}, \beta_0 = 0$$

$$\text{pvalue} = 2 \cdot P(T \leq -|t_{\hat{\beta}}|) \sim t_{n-p}$$

$$df = n - p$$

$$R_{adj}^2 = 1 - \frac{(1 - R^2)(n-1)}{n-p}$$

$$F = \frac{\left(\sum_{i=1}^n (y_i - \bar{y})^2 - \sum_{i=1}^n e_i^2 \right) / (p-1)}{\sum_{i=1}^n e_i^2 / (n-p)} = \frac{(TSS - RSS) / (p-1)}{RSS / (n-p)}$$

$$df_{num} = p - 1$$

$$df_{denom} = n - p$$

$$P(F > f) \sim F_{df_{num}, df_{denom}}$$

Predicting the Future

... with confidence I can predict what happened yesterday ...

```
# Values to use for prediction
```

```
future_car = data.frame(wt = 0.5, qsec = 12)
```

```
# Make prediction
```

```
y_hat = predict(model_fit, new_data = future_car)
```

```
head(y_hat)
```

#	Mazda RX4	Mazda RX4 Wag	Datsun 710
#	21.81511	21.04822	25.32728
#	Hornet 4 Drive	Hornet Sportabout	Valiant
#	21.58057	18.19611	21.06859

Factors

Previously

Hierarchy of Data Types

*

Variable Data

* **raw** type is missing

Numerical

Categorical

Continuous

Discrete

character

"toad", "got pie?", "orange", "stat 385"
Strings

complex

$3 + 4i$, $4.2 - i$, $3.6i$, $99 + 0i$
Real and imaginary

integer

6, 12, 0, -4
Whole Numbers

Nominal

Ordinal

numeric

3.14, 10.5, 0.0, -4.8
Decimals

logical

TRUE (1), FALSE (0)
Boolean values

factor

"AB", "A", "B", "O"
Blood Type

ordered factor

"A", "B", "C", "D", "F"
Letter grades

Augmented

Verifying Class in R

subject_heights

id	sex	height
1	M	6.1
2	F	5.5
3	F	5.2
...
55	M	5.9

```
id = c(1, 2, 3, 55)
```

```
class(id)  
# [1] "numeric"
```

```
sex = c("M", "F", "F", "M")
```

```
class(sex)  
# [1] "character"
```


```
height = c(6.1, 5.5, 5.2, 5.9)
```

```
class(height)  
# [1] "height"
```

Class Difference


```
# Default way
# Treat as factors
subject_heights = data.frame(
  id      = c(1, 2, 3, 55),
  sex     = c("M", "F", "F", "M"),
  height  = c(6.1, 5.5, 5.2, 5.9),
  stringsAsFactors = TRUE
)
summary(subject_heights)
```

	id	sex	height
Min.	: 1.00	F:2	Min. :5.200
1st Qu.:	1.75	M:2	1st Qu.:5.425
Median :	2.50		Median :5.700
Mean :	15.25		Mean :5.675
3rd Qu.:	16.00		3rd Qu.:5.950
Max.	:55.00		Max. :6.100



```
# Opt out of factors
# Treat as characters
subject_heights_nofct = data.frame(
  id      = c(1, 2, 3, 55),
  sex     = c("M", "F", "F", "M"),
  height  = c(6.1, 5.5, 5.2, 5.9),
  stringsAsFactors = FALSE
)
summary(subject_heights_nofct)
```

	id	sex	height
Min.	: 1.00	Length:4	Min. :5.200
1st Qu.:	1.75	Class :character	1st Qu.:5.425
Median :	2.50	Mode :character	Median :5.700
Mean :	15.25		Mean :5.675
3rd Qu.:	16.00		3rd Qu.:5.950
Max.	:55.00		Max. :6.100



stringsAsFactors: An unauthorized biography

👤 Roger Peng 📅 2015/07/24

Recently, I was listening in on the conversation of some colleagues who were discussing a bug in their R code. The bug was ultimately traced back to the well-known phenomenon that functions like `'read.table()'` and `'read.csv()'` in R convert columns that are detected to be character/strings to be factor variables. This lead to the spontaneous outcry from one colleague of

Why does `stringsAsFactors` not default to `FALSE`????

The argument `'stringsAsFactors'` is an argument to the `'data.frame()'` function in R. It is a logical that indicates whether strings in a data frame should be treated as factor variables or as just plain strings. The argument also appears in `'read.table()'` and related functions because of the role these functions play in reading in table data and converting them to data frames. By default, `'stringsAsFactors'` is set to `TRUE`.

This argument dates back to May 20, 2006 when it was originally introduced into R as the `'charToFactor'` argument to `'data.frame()'`. Soon afterwards, on May 24, 2006, it was changed to `'stringsAsFactors'` to be compatible with S-PLUS by request from Bill Dunlap.

Most people I talk to today who use R are completely befuddled by the fact that `'stringsAsFactors'` is set to `TRUE` by default. First of all, it should be noted that before the `'stringsAsFactors'` argument even existed, the behavior of R was to coerce all character strings to be factors in a data frame. If you didn't want this behavior, you had to manually coerce each column to be character.

So here's the story:

In the old days, when R was primarily being used by statisticians and statistical types, this setting strings to be factors made total sense. In most tabular data, if there were a column of the table that was non-numeric, it almost certainly encoded a categorical variable. Think sex (male/female), country (U.S./other), region (east/west), etc. In R, categorical variables are represented by `'factor'` vectors and so character columns got converted factor.

Why do we need factor variables to begin with? Because of modeling functions like `'lm()'` and `'glm()'`. Modeling functions need to treat expand categorical variables into individual dummy variables, so that a categorical variable with 5 levels will be expanded into 4 different columns in your modeling matrix. There's no way for R to know it should do this unless it has some extra information in the form of the factor class. From this point of view, setting `'stringsAsFactors = TRUE'` when reading in tabular data makes total sense. If the data is just going to go into a regression model, then R is doing the right thing.

<https://simplystatistics.org/2015/07/24/stringsasfactors-an-unauthorized-biography/>

Just say **YES** to
stringsAsFactors = FALSE

... or feel the wraith of factors ...

Factors

Codifying character values as numbers

```
sex = c("M", "F", "F", "M") # Character vector
```

```
sex
```

```
# [1] "M" "F" "F" "M"
```


```
sex_factor = factor(sex) # Convert from character to factor
```

```
# [1] M F F M # No quotations around character!
```

```
# Levels: F M # Lists the character mapped to numbers.
```

```
as.numeric(sex_factor) # Access internal level / number representation
```

```
# [1] 2 1 1 2
```



* Storing data as factors instead of character data on disk will result in lower file sizes.

** Computationally, there are advantages by transforming character values to factors.

Usefulness of Factors

- Creating a **design matrix** for *linear regression* as they provide a *codified dummy variable* structure. e.g. FALSE (0) or TRUE (1)

sex	height
"M"	6.1
"F"	5.5
"F"	5.2
...	...
"M"	5.9

subject_data



sex_F	height
0	6.1
1	5.5
1	5.2
...	...
0	5.9

Design Matrix

- Poor if character values (e.g. levels) need to be manipulated or there are too many unique values (e.g. text messages on a phone.)

Your Turn

create the design matrix for where students sit in class

dist	side		dist_back	side_middle	side_right
"back"	"right"				
"back"	"left"				
"front"	"middle"	→			
"back"	"middle"				
"front"	"right"				
"front"	"left"				

Why are *side_left* or *dist_front* not included as variables?

Limitations of Factors

... no math support ...

```
x = c(3L, -1L, 22L, 9L, 0L, 22L, 9L) # Create Integer Vector
```

```
my_factor = as.factor(x) # Cast integer to factor
```

```
my_factor + 10 # Error in an unexpected way
```

```
# Warning in Ops.factor(my_factor, 10) : '+' not meaningful for factors
```

```
# [1] NA NA NA NA NA NA NA
```

```
min(my_factor) # Show stopping error
```

```
# Error in Summary.factor(c(3L, 1L, 5L, 4L, 2L, 5L, 4L), na.rm = FALSE) :
```

```
# 'min' not meaningful for factors
```


Fun with Factors

... viewing levels, changing values, and ...

```
levels(my_factor)  
# [1] "-1" "0"  "3"  "9"  "22"
```

List of Levels

```
my_factor[1] = 9  
# [1] 9 -1 22 9 0 22 9  
# Levels: -1 0 3 9 22
```

Modify with a pre-existing level

```
my_factor[2] = 18  
# Warning in `[<-.factor`(`*tmp*`, 2, value = 18) :  
# invalid factor level, NA generated
```

Error if level isn't present already

Recovering Values from a Factor

Translating a factor's levels to an atomic vector

Extract out names of levels by element-wise position.

For details, see ?`[.factor`

x_char = levels(my_factor)[my_factor] # Levels treated as positions

x_char

[1] "3" "-1" "22" "9" "0" "22" "9" # Character vector extracted.

Coerce to the appropriate type. (Integer for this example...)

x_val = **as.integer**(x_char)

x_val

[1] 3 -1 22 9 0 22 9 # Integer vector returned so data is recovered!

my_factor		levels
"3"	↔	3
"-1"	↔	1
"22"	↔	5
"9"	↔	4
"0"		2
"22"		5
"9"		4

Factor Structure

levels(my_factor)	
id	my_factor
1	"-1"
2	"0"
3	"3"
4	"9"
5	"22"

Mapping

levels
3
1
5
4
2
5
4

Level
Position

x_char	x_val
"3"	3
"-1"	-1
"22"	22
"9"	9
"0"	0
"22"	22
"9"	9

Character ➔ Integer
Explicit Coercion

Ordered Factors

... making factors have precedence ...

```
yields = c("hi", "low", "med", "low", "med", "low") # Vector of Character Values
```

```
yields_fct = factor(yields) # Create factor  
# [1] hi low med low med low  
# Levels: hi low med
```

```
yields_ordered = factor(yields, ordered = TRUE) # Add ordered component  
# [1] hi med low med low med low  
# Levels: hi < low < med
```

```
# Correct ordering from low to high  
yields_fixed_order = factor(yields, levels = c("low", "med", "hi"), ordered = TRUE)  
# [1] hi med low med low med low  
# Levels: low < med < hi
```

Your Turn

Determine whether the following should be either
a factor or an ordered factor

Months: (Jan, Feb, ... , Nov, Dec)

Colors: (red, orange, ... , black, green)

Alphabet: (a, b, ... , y, z)

Recap

- **SLR and MLR**

- Estimating a linear regression with 2 parameters vs. p parameters
- Underlying design matrix construction

- **Factors**

- Provide a mapping of levels to indicator variables inside the design matrix.

This work is licensed under the
Creative Commons
Attribution-NonCommercial-
ShareAlike 4.0 International
License

